



中山大學

SUN YAT-SEN UNIVERSITY

Lecture 15.

Neural Networks

Pattern Recognition and Computer Vision

Guanbin Li,


School of Computer Science and Engineering, Sun Yat-Sen University

扫码签到



Perceptron


一个神经元通常具有多个**树突**，主要用来接受传入信息；而**轴突**只有一条，轴突尾端有许多轴突末梢可以给其他多个神经元传递信息。轴突末梢跟其他神经元的树突产生连接，从而传递信号。这个连接的位置在生物学上叫做“**突触**”。




输入：权重 W , 初始值 X

Score function: $s = WX$


输出: $Y = f(WX)$, 其中 f 是激活函数



Perceptron



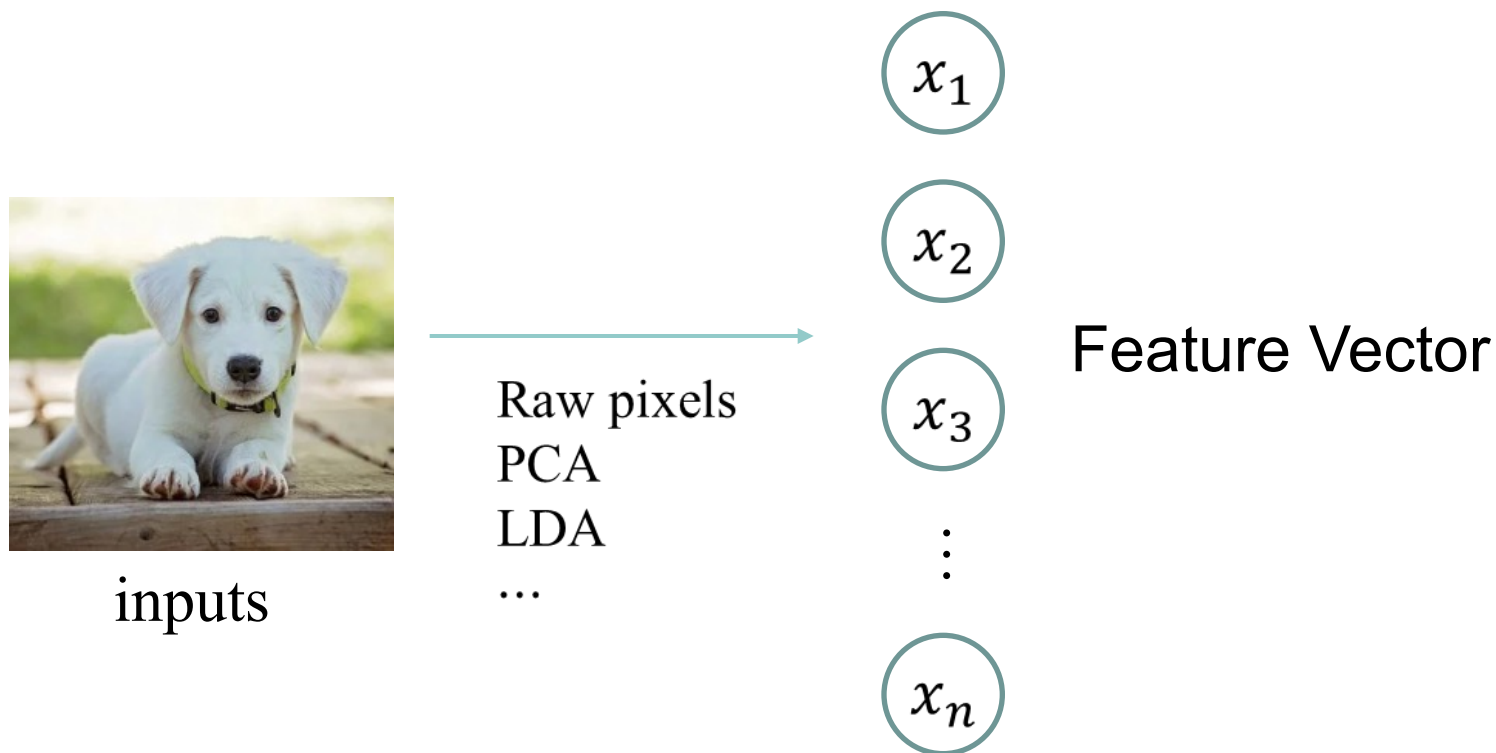
Perceptron



What are the inputs?


Perceptron

- For image classification



Perceptron

- For image classification



How do we handle multi-class classification task?


Add more perceptrons!

Perceptron

- Two perceptrons



inputs




Perceptron

- Three perceptrons



inputs



Linear classifier is a set of perceptrons

What we will learn today?


- Perceptron
- **Linear classifier**
- Loss function
- Gradient descent and backpropagation
- Neural networks

Linear classifier

- Input layer



inputs




Input layer

Linear classifier

- Output layer

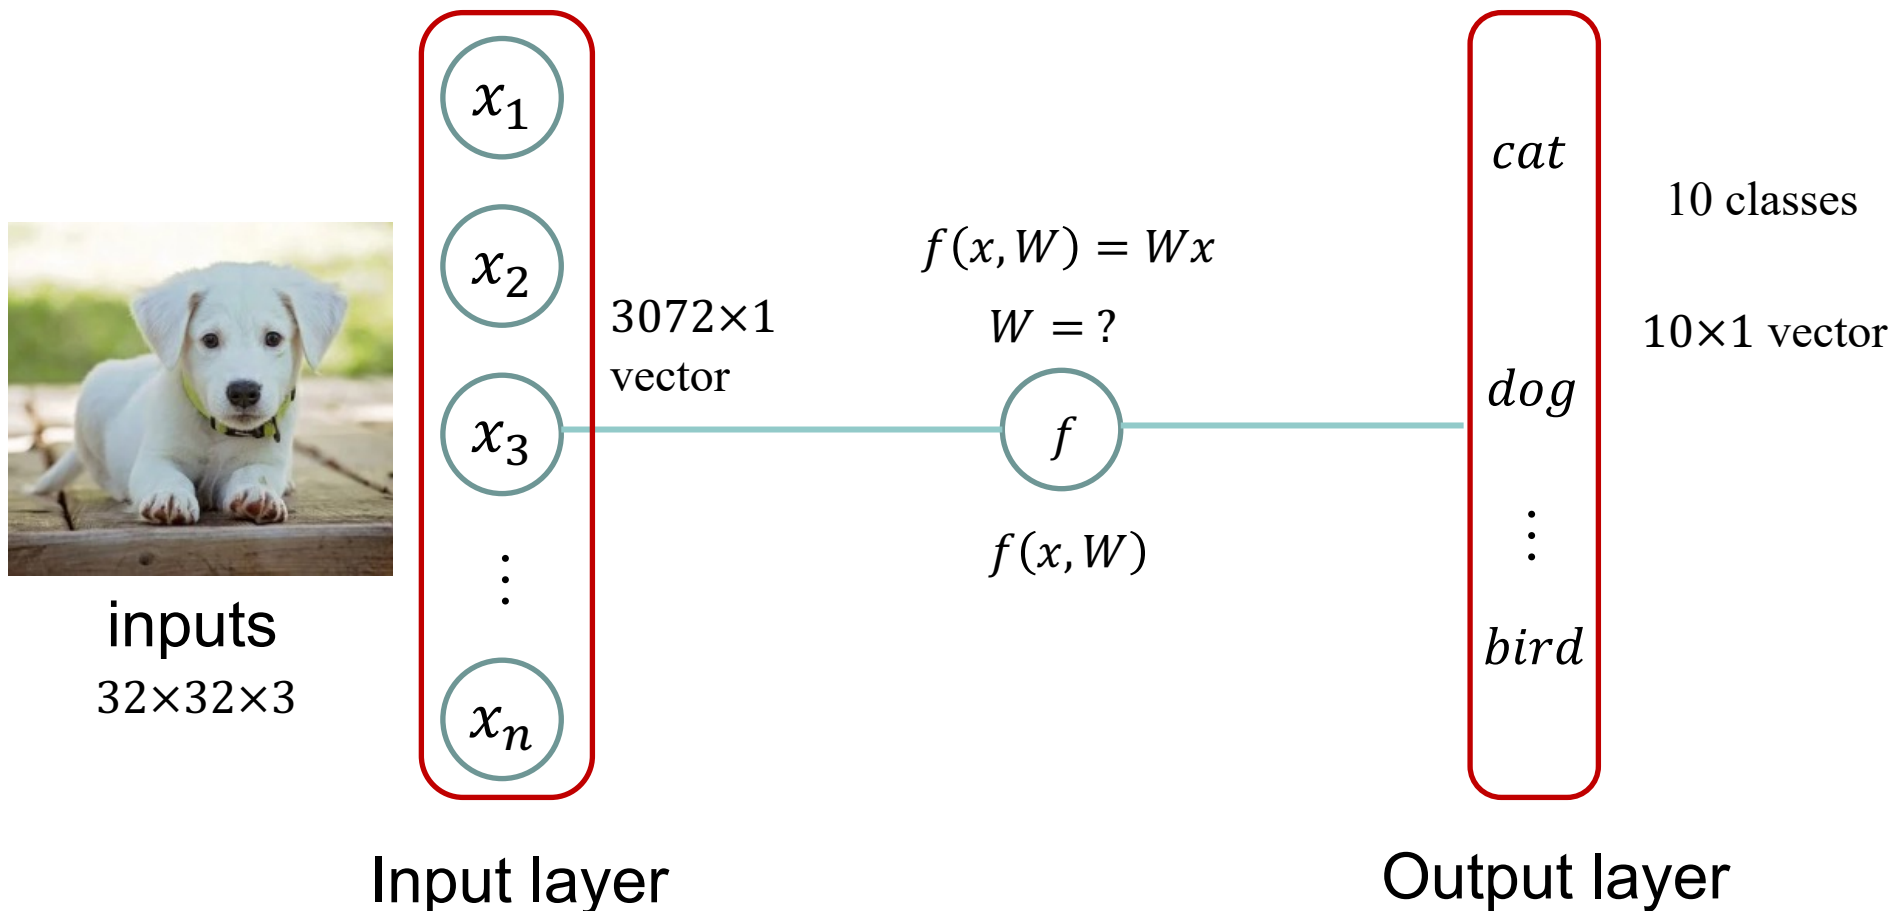


inputs



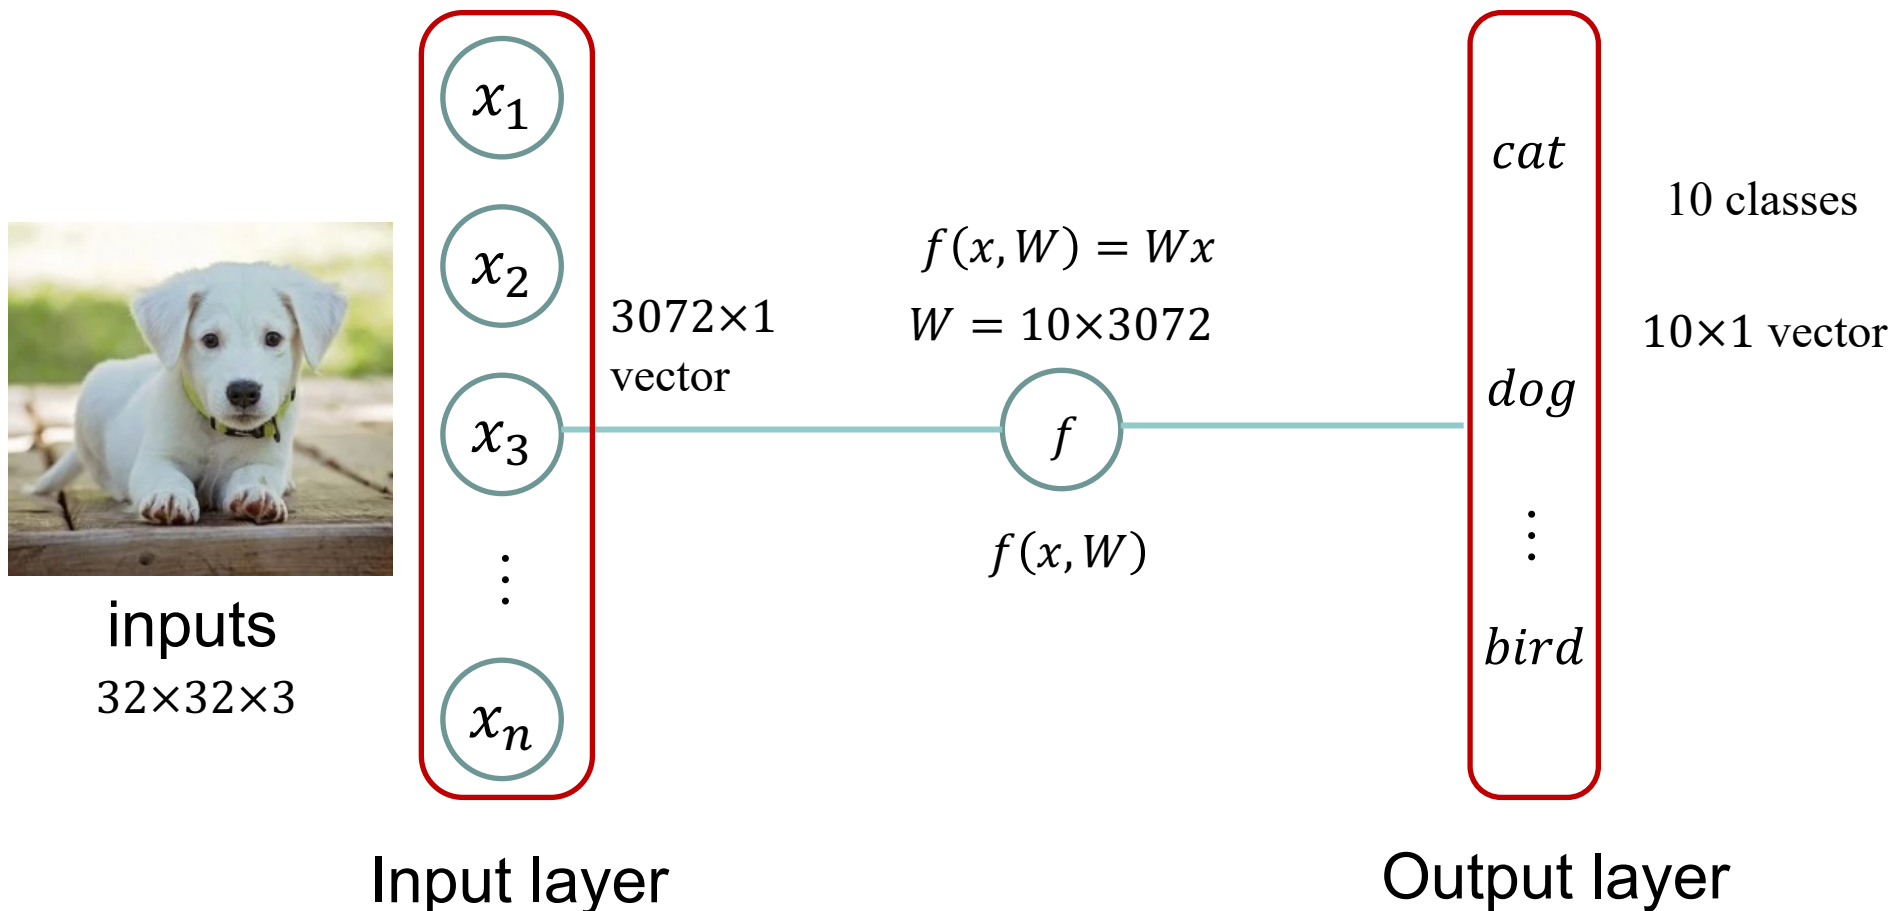
Linear classifier

- Mathematical formulation




Linear classifier

- Mathematical formulation




Linear classifier

- Classification task



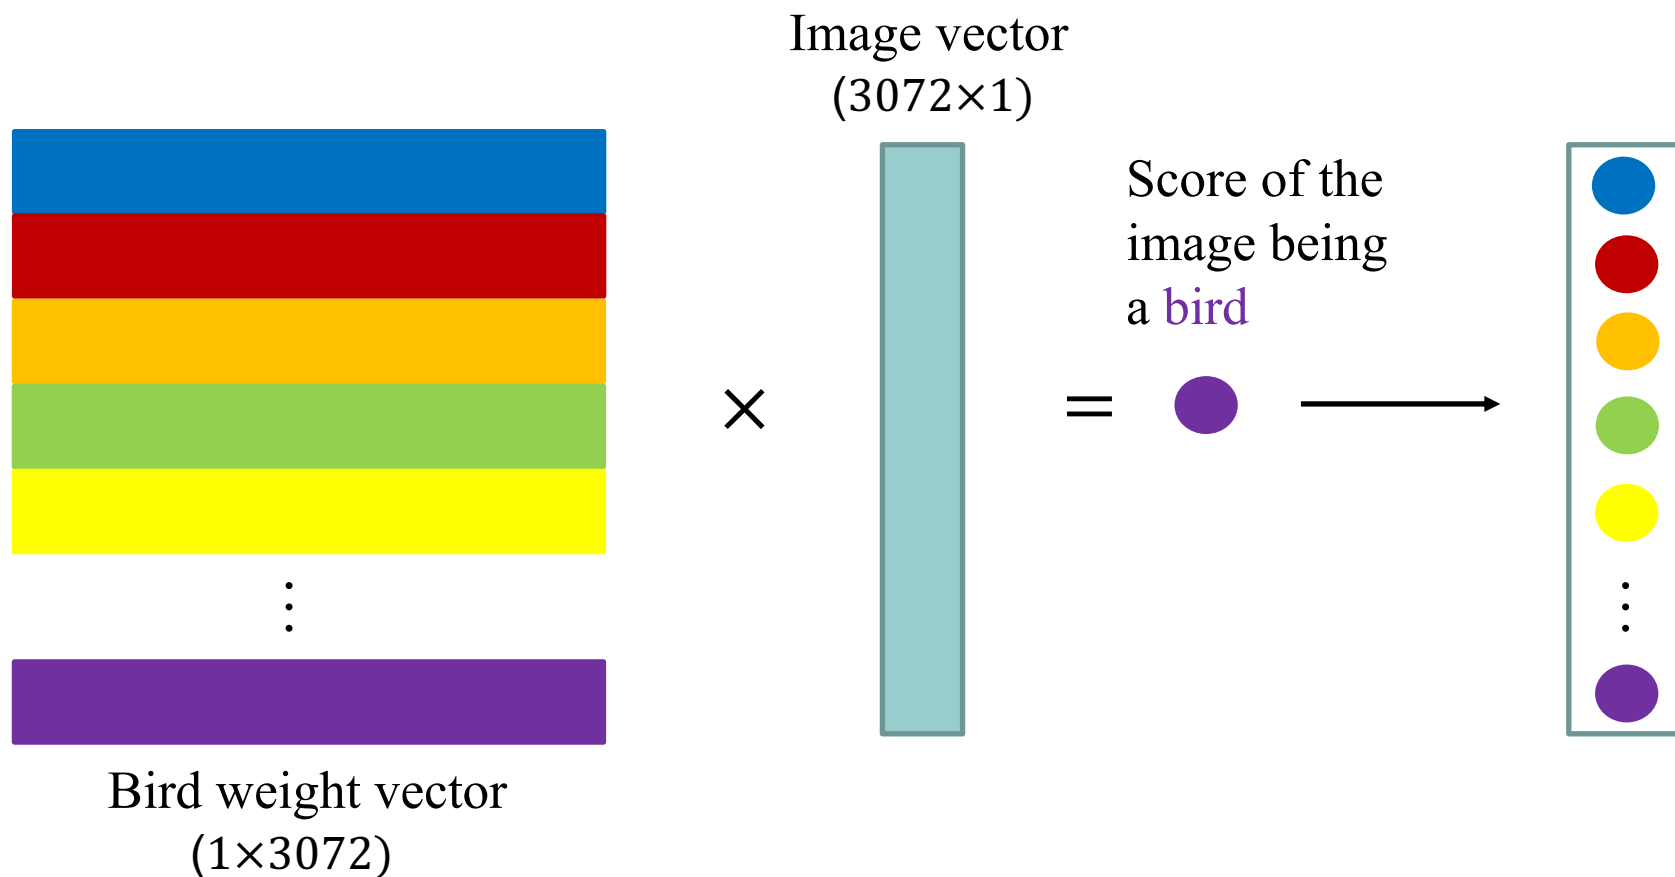
Linear classifier

- Classification task




Linear classifier

- Classification task



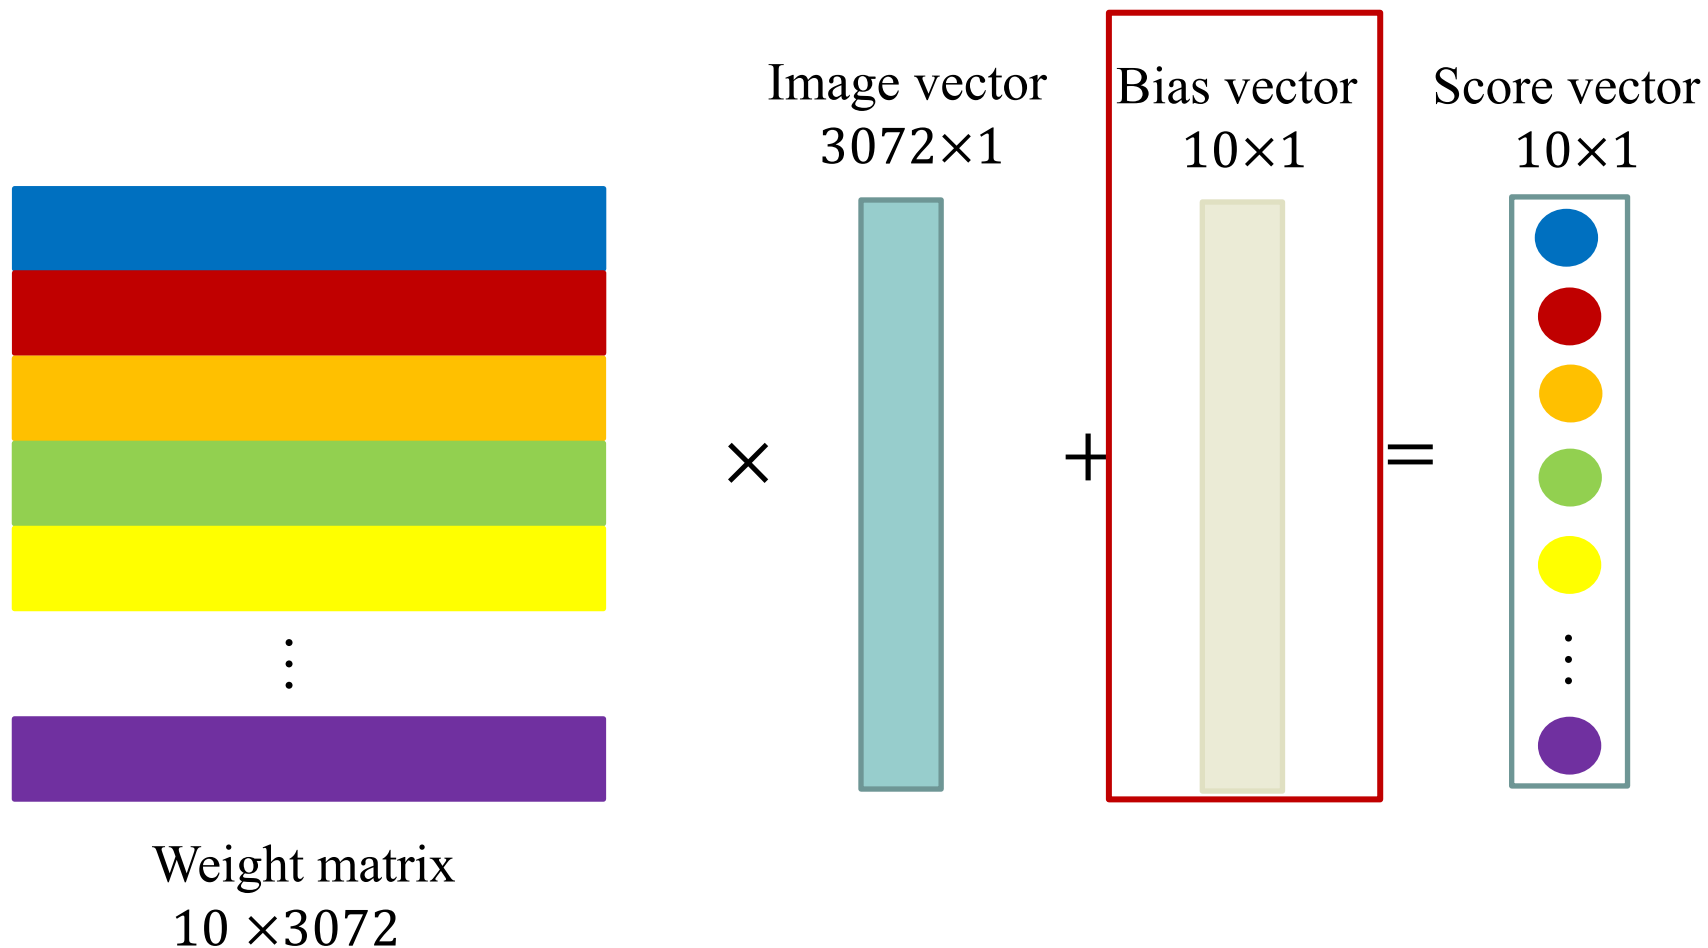
Linear classifier

- Classification task



Linear classifier


- Classification task




Linear classifier

- Bias can enhance the learning ability of the network

$$y = Wx$$

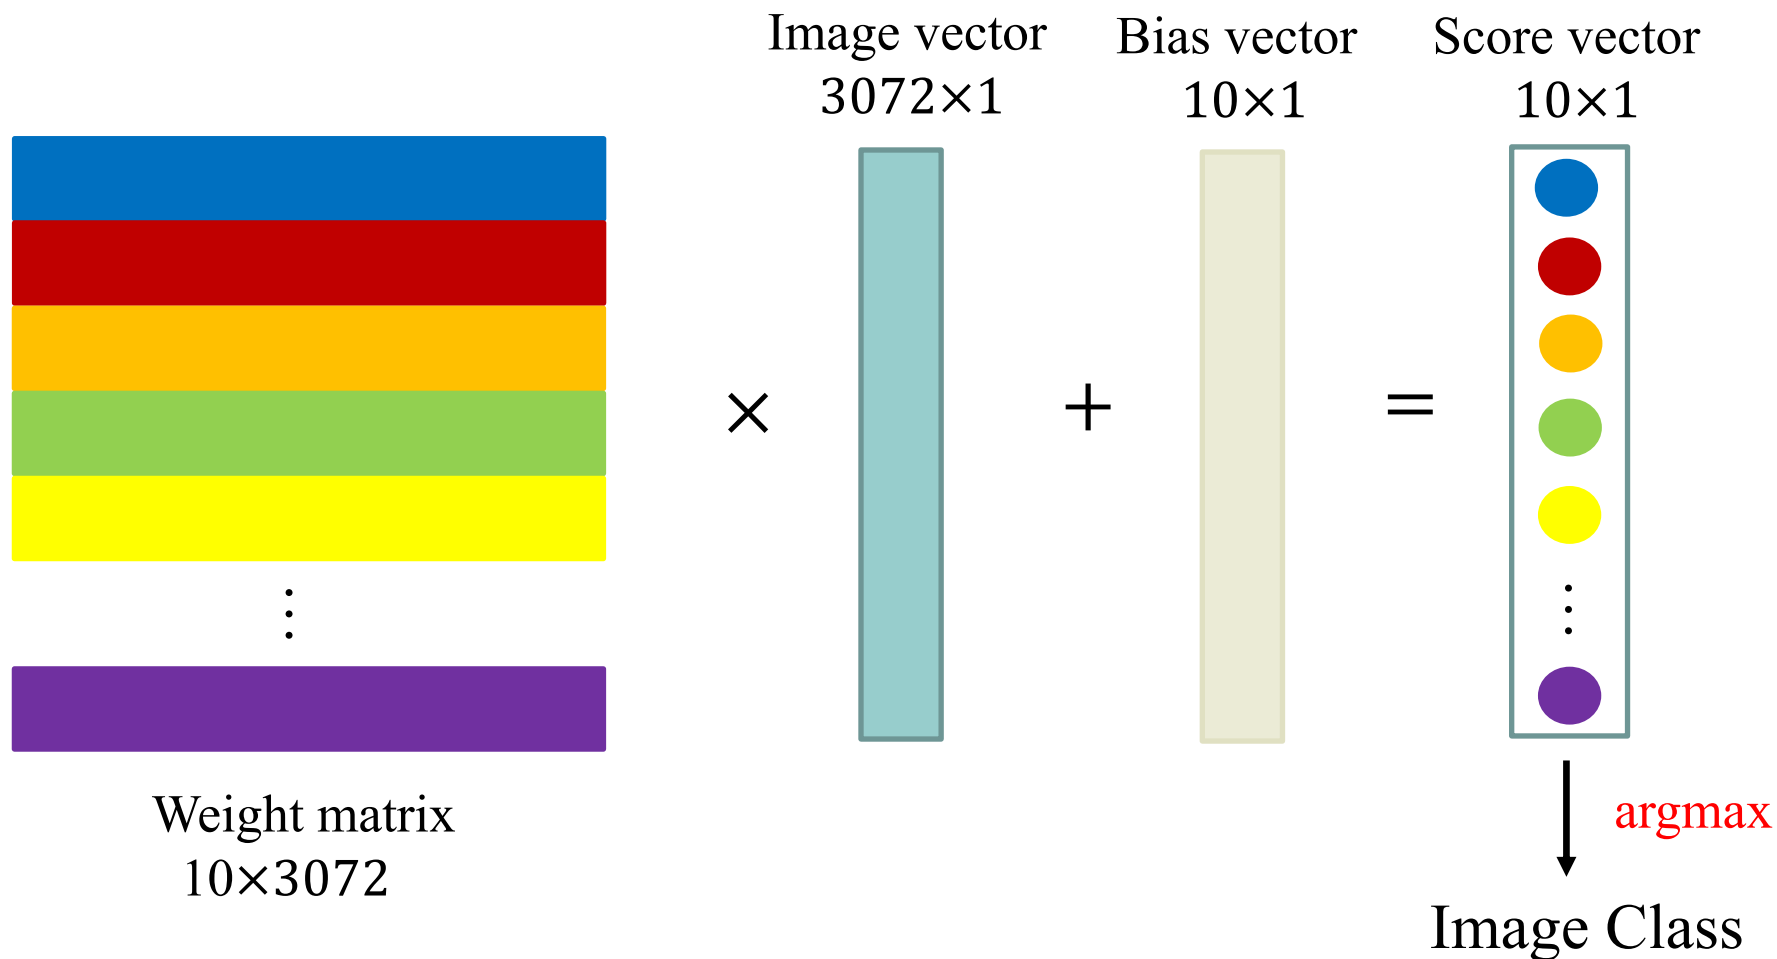


$$y = Wx + b$$



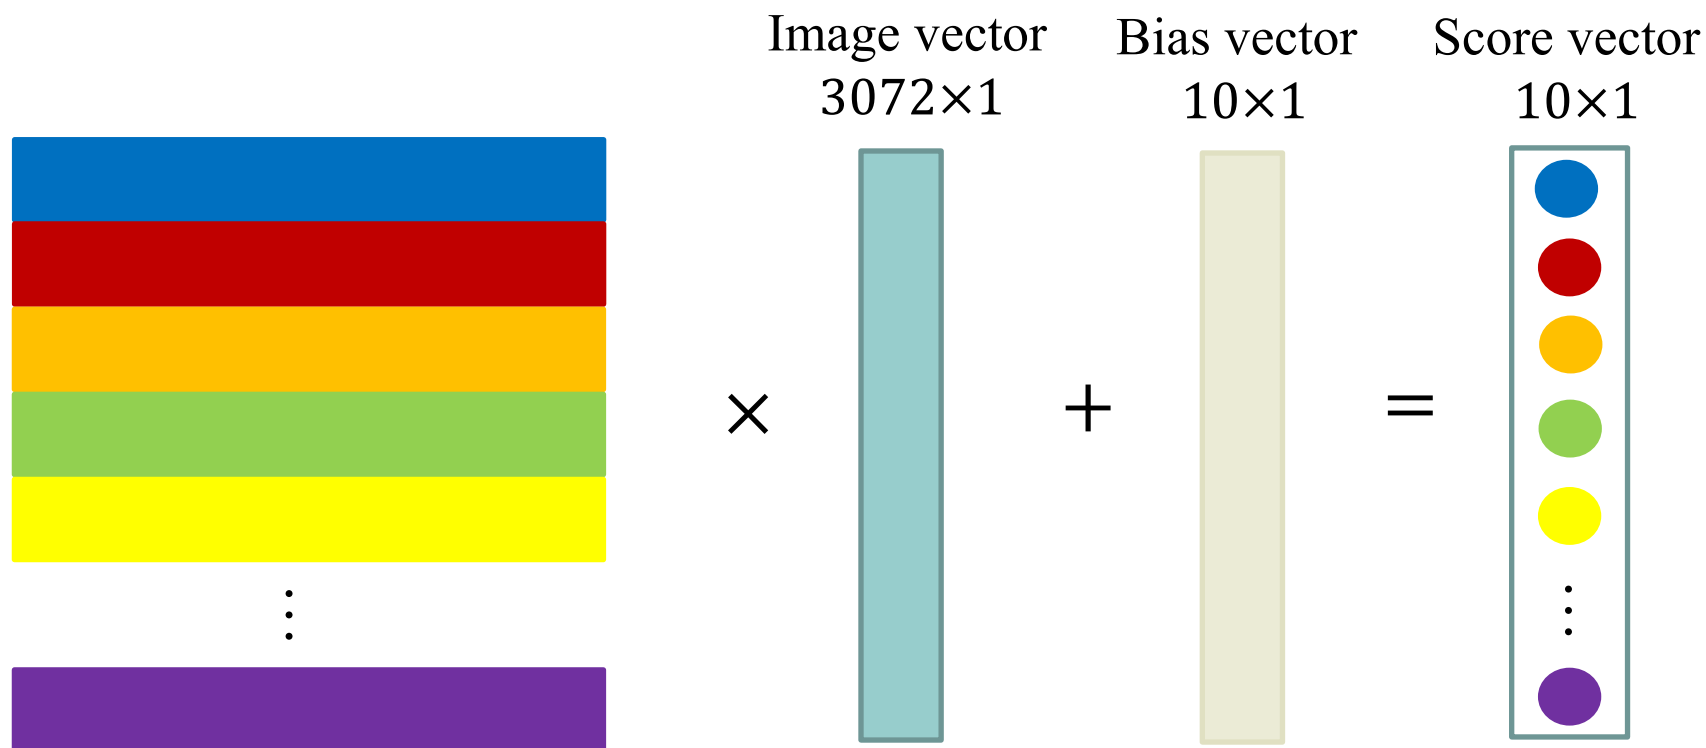
Linear classifier

- Classification task



Linear classifier

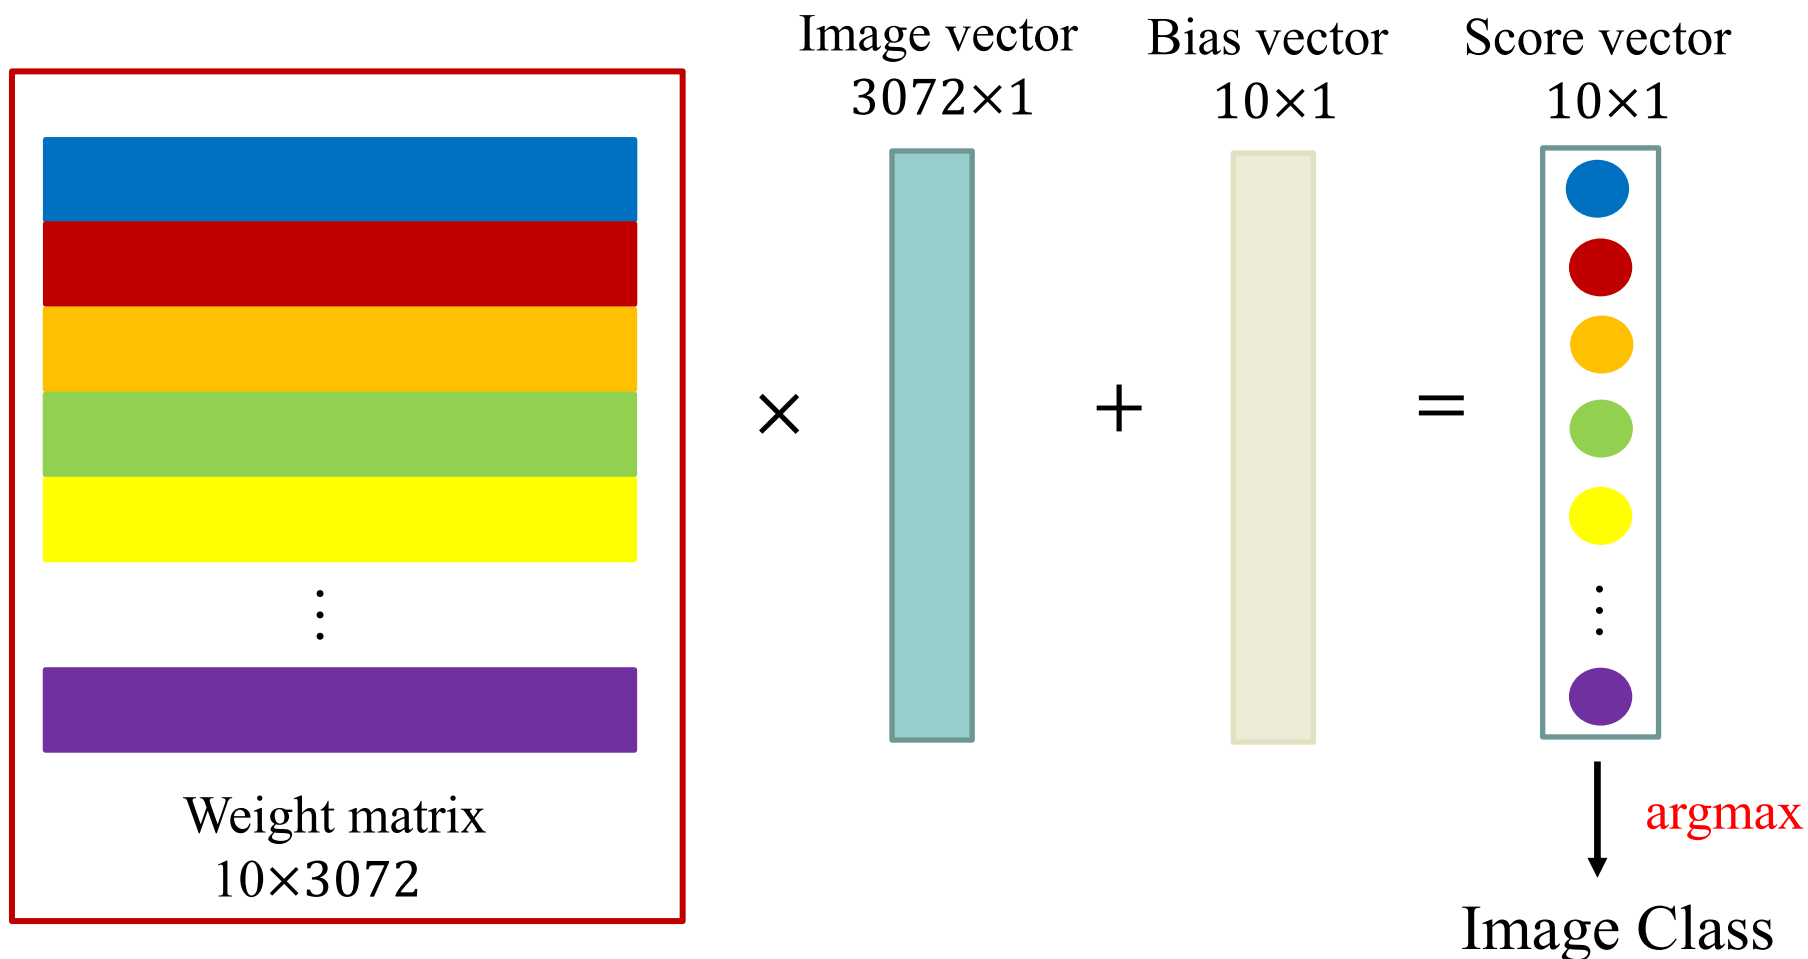
- Classification task



How do the classifier get the highest score for the right class?

Linear classifier


- The key of classification task is the **weight matrix**



Linear classifier

- The key of classification task is the weight matrix

The weight matrix can be obtained by **training!**



What we will learn today?

- Perceptron
- Linear classifier
- **Loss function**
- Gradient descent and backpropagation
- Neural networks

Loss function

- Loss function determines how good our classifier

Given some training examples:

$$\{(x_1, y_1), (x_2, y_2), (x_3, y_3), \dots (x_n, y_n)\}$$

x_i the i th image, y_i is the corresponding integer label

(e.g. 0 for dog, 1 for cat ...)

and our classifier: $\hat{y} = Wx$

Loss of one example is determined as $L_i(y_i, \hat{y}_i)$

when the classifier predicts correctly ($\hat{y}_i = y_i$), the loss is low

when the classifier makes mistakes ($\hat{y}_i \neq y_i$), the loss is high

Loss function

- Loss function determines how good our classifier

Given some training examples:

$$\{(x_1, y_1), (x_2, y_2), (x_3, y_3), \dots (x_n, y_n)\}$$

x_i the i th image, y_i is the corresponding integer label
(e.g. 0 for dog, 1 for cat ...)

and our classifier: $\hat{y} = Wx$

Loss over the dataset is the average of loss over examples:

$$Loss = \frac{1}{N} \sum_i^N L_i(y_i, \hat{y}_i)$$

Loss function

- Loss function is the key to find suitable W

Specifically, we need to find W such that:

$$\min_w \text{Loss}(y, \hat{y})$$

y is the true labels, \hat{y} is the model predicted labels.

Loss function

- Some popular loss functions

L1 Loss

$$L_i(y_i, \hat{y}_i) = |y_i - \hat{y}_i|$$

Squared Error Loss(L2 Loss)

$$L_i(y_i, \hat{y}_i) = (y_i - \hat{y}_i)^2$$

Zero-One Loss

$$L_i(y_i, \hat{y}_i) = 1 \|y_i \neq \hat{y}_i\|$$

Hinge Loss

$$L_i(y_i, \hat{y}_i) = \max(0, 1 - y_i \hat{y}_i)$$

Loss function

- Softmax Loss (multi-class classification task)
- Softmax function allows us to treat the outputs of a model as **probabilities** for each class.
- Common way of measuring distance between **probability distributions** is Kullback-Leibler (KL) divergence.

$$D_{KL} = \sum_y P(y) \log \frac{P(y)}{Q(y)}$$

P is the ground truth distribution and

Q is the output score distribution

Loss function

- KL divergence

$$D_{KL} = \sum_y P(y) \log \frac{P(y)}{Q(y)}$$

In our case, P is only non-zero for the correct class.

For example, consider the case where we only have 3 classes:



1	dog
0	cat
0	bird

Ground Truth

Loss function

- KL divergence

$$D_{KL} = \sum_y P(y) \log \frac{P(y)}{Q(y)}$$

$$= -\log Q(y) \text{ when } y = \text{dog}$$

$$= -\log \text{Prob}(f(x_i, W) = y_i)$$



1	dog
0	cat
0	bird

Ground Truth

Loss function

- KL divergence

$$L_i = -\log \text{Prob}(f(x_i, W) == y_i)$$

Our linear classifier:

$$\hat{y} = wx$$

There is no limits on the output space.

Meaning that the model can generate outputs >1 or <0 .



3.2
5.1
-1.7

Model Outputs

1	dog
0	cat
0	bird

Ground Truth

Loss function

- Softmax function

$$L_i = -\log \text{Prob}(f(x_i, W) == y_i)$$

We need to **convert the outputs into probability ranges [0,1]**.



3.2
5.1
-1.7

Model Outputs

1	dog
0	cat
0	bird

Ground Truth

Loss function

- Softmax function

$$L_i = -\log \text{Prob}(f(x_i, W) == y_i)$$

We need to convert the outputs into probability ranges $[0,1]$.

Solution: Softmax: $\text{Prob}(f(x_i, W) == k) = \frac{e^{\hat{y}^k}}{\sum_j e^{\hat{y}^j}}$



3.2
5.1
-1.7

Model Outputs

1	dog
0	cat
0	bird

Ground Truth

Loss function

- Softmax function


$$L_i = -\log \text{Prob}(f(x_i, W) == y_i)$$

We need to convert the outputs into probability ranges $[0,1]$.

Solution: Softmax: $\text{Prob}(f(x_i, W) == k) = \frac{e^{\hat{y}^k}}{\sum_j e^{\hat{y}^j}}$



Model Outputs



Ground Truth


Loss function

- Softmax function

$$L_i = -\log \text{Prob}(f(x_i, W) == y_i)$$

In this case, loss is: $L_i = -\log(0.13) = 2.04$

The dog probability closes to 1, the loss closes to 0.




What we will learn today?

- Perceptron
- Linear classifier
- Loss function
- Gradient descent and backpropagation
- Neural networks


Gradient descent

- Visualization: Minimizing loss




Gradient descent

- Visualization: Minimizing loss




Gradient descent

- Visualization: Minimizing loss




Gradient descent

- Visualization: Minimizing loss




Gradient descent

- Visualization: Minimizing loss




Gradient descent

- Visualization: Minimizing loss



Gradient descent

- Visualization: Minimizing loss



Gradient descent

- Gradient Descent Pseudocode

for $_$ in $\{0, \dots, num_epochs\}$:

$L = 0$

for x_i, y_i in data:

$$\hat{y}_i = f(x_i, W)$$

$$L += L_i(y_i, \hat{y}_i)$$

$$\frac{dL}{dW} = ???$$

$$W := W - \alpha \frac{dL}{dW}$$

Learning rate

Gradient descent

- Partial derivative of loss to update weights

Given training data point (x, y) , the linear classifier formula is: $\hat{y}_i = Wx$

Let's assume that the correct label is class k , implying $y = k$

$$Loss = L(\hat{y}, y) = -\log \frac{e^{\hat{y}^k}}{\sum_j e^{\hat{y}^j}} \quad (\text{Softmax loss})$$

$$= -\hat{y}_k + \log \sum_j e^{\hat{y}^j}$$

Calculating the loss $\frac{dL}{dW}$ is hard mathematically. But we can use the **chain rule** to make it simpler:

$$\frac{dL}{dW} = \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dW}$$

Gradient descent

- Partial derivative of loss to update weights

Given training data point (x, y) , the linear classifier formula is: $\hat{y}_i = Wx$

Let's assume that the correct label is class k , implying $y = k$

$$Loss = L(\hat{y}, y) = -\log \frac{e^{\hat{y}^k}}{\sum_j e^{\hat{y}^j}} \quad (\text{Softmax loss})$$

$$= -\hat{y}_k + \log \sum_j e^{\hat{y}^j}$$

Calculating the loss $\frac{dL}{dW}$ is hard mathematically. But we can use the chain rule to make it simpler:

$$\frac{dL}{dW} = \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dW}$$

We know that $\frac{d\hat{y}}{dW} = x$, but what about $\frac{dL}{d\hat{y}}$?

Gradient descent

- Partial derivative of loss to update weights

$$L = -\hat{y}_k + \log \sum_j e^{\hat{y}^j}$$

To calculate $\frac{dL}{d\hat{y}}$, we need to consider two cases:

Case1:

$$\frac{dL}{d\hat{y}_k} = -1 + \frac{e^{\hat{y}^k}}{\sum_j e^{\hat{y}^j}}$$

Case2:

$$\frac{dL}{d\hat{y}_{l \neq k}} = \frac{e^{\hat{y}^l}}{\sum_j e^{\hat{y}^j}}$$

Gradient descent

- Partial derivative of loss to update weights

Put it all together:

$$\frac{dL}{dW} = \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dW}$$

$$\frac{dL}{dW} = \begin{pmatrix} \frac{e^{\hat{y}0}}{\sum_j e^{\hat{y}j}} \\ \dots \\ -1 + \frac{e^{\hat{y}k}}{\sum_j e^{\hat{y}j}} \\ \dots \\ \frac{e^{\hat{y}3071}}{\sum_j e^{\hat{y}j}} \end{pmatrix} x$$

Gradient descent

- Gradient Descent Pseudocode

for $_$ in $\{0, \dots, num_epochs\}$:

$L = 0$

for x_i, y_i in data:

$$\hat{y}_i = f(x_i, W)$$

$$L += L_i(y_i, \hat{y}_i)$$

$$\frac{dL}{dW} = \textit{We know how to calculate this now!}$$

$$W := W - \alpha \frac{dL}{dW}$$

Gradient descent

- Gradient Descent Pseudocode

for $_$ in $\{0, \dots, num_epochs\}$:

$L = 0$

for x_i, y_i in data:

$$\hat{y}_i = f(x_i, W)$$

$$L += L_i(y_i, \hat{y}_i)$$


$$\frac{dL}{dW} = \textit{We know how to calculate this now!}$$

$$W := W - \alpha \frac{dL}{dW}$$

After num_epochs , W is well suited to the classification task

Backprop

- Backprop – another way of computing gradients
 - visualize the computation as a graph
 - Compute the **forward pass** to calculate the loss.
 - Compute all **gradients** for each computation **backwards**



$$\hat{y} = Wx$$

$$L = \text{Loss}(\hat{y}, y)$$


$$\frac{dL}{dW} = \frac{dL}{d\hat{y}} \frac{d\hat{y}}{dW}$$

What we will learn today?

- Perceptron
- Linear classifier
- Loss function
- Gradient descent and backpropagation
- Neural networks

Neural networks


- Features sometimes might not be linearly separable



What should we do in this case?

Neural networks

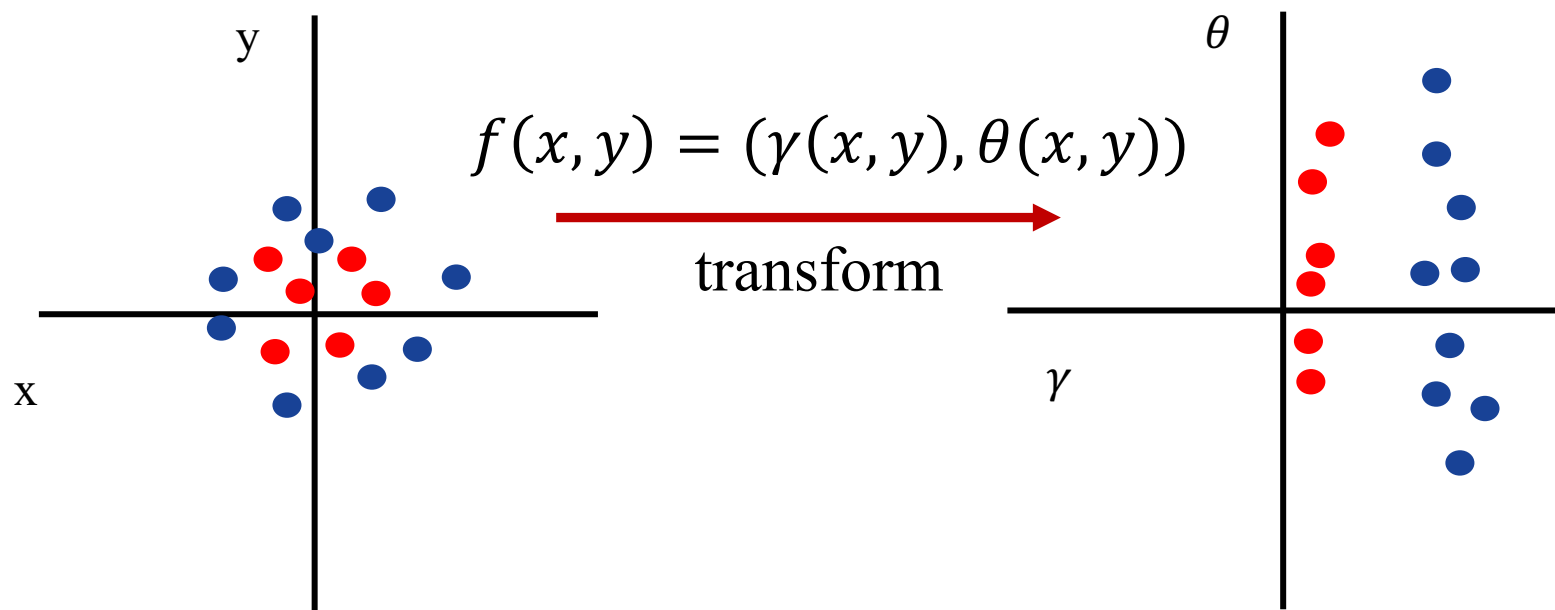
- Features sometimes might not be linearly separable



We can transform it to be linearly separable!

Neural networks

- Features sometimes might not be linearly separable




Neural networks

- Recall: Our linear classifier



inputs




Neural networks

- Transform the features by adding another layer



inputs



Neural networks

- 2-layer network: mathematical formula

linear classifier:

$$y = Wx$$

2-layer network:

$$y = W_2(\max(0, W_1x))$$

We know the size of $x = 3072 \times 1$ and $y = 10 \times 1$

So:

$$W_1 = h \times 3072 \text{ and } W_2 = 10 \times h$$

h is a new hyperparameter!

Neural networks

- 2-layer network: mathematical formula

linear classifier:

$$y = Wx$$

2-layer network:

$$y = W_2(\max(0, W_1x))$$

Activation function

We know the size of $x = 3072 \times 1$ and $y = 10 \times 1$

So:

$$W_1 = h \times 3072 \text{ and } W_2 = 10 \times h$$

h is a new hyperparameter!

Neural networks

- 2-layer network: mathematical formula


linear classifier:

$$y = Wx$$

2-layer network:

$$y = W_2(\max(0, W_1x))$$

Activation function



Why is the activation function necessary?

If we remove it, the neural network turns to be a linear classifier

$$y = W_2W_1x = Wx$$


Neural networks

- Activation function

It allows models to learn complex transformations for features.
Choosing the right activation function is important!


a. Sigmoid (binary classification)

$$\delta(x) = \frac{1}{1 + e^{-x}}$$



b. tanh

$$\tanh(x)$$




Neural networks

- Activation function

It allows models to learn complex transformations for features.
Choosing the right activation function is important!


c. **ReLU (Rectified Linear Unit, ReLU)**

$$\max(0, x)$$




d. **Leaky ReLU**

$$\max(0.1x, x)$$



Neural networks

- Multi-layers Neural networks



What we have learned today?

- Perceptron
- Linear classifier
- Loss function
- Gradient descent and backpropagation
- Neural networks

What we have learned today?

- Perceptron----Inspiration from Biology
- Linear classifier----A set of perceptrons
- Loss function----Determines how good our classifier
- Gradient descent and backpropagation----Training
- Neural networks---Handle complex feature distribution

Questions?